



A Valid-response counts by variant**B Loss does not predict valid behavior**

Valid Ready-Set-Go responses per seed across learning rules. (A) BPTT and RFLO produce valid responses reliably; PC fails at 20 inference steps and partially recovers at 100. (B) Best training loss (log scale) does not predict valid behavior. Points are individual seeds; bars in A are per-variant means.